


Data Manipulation I

John Rios

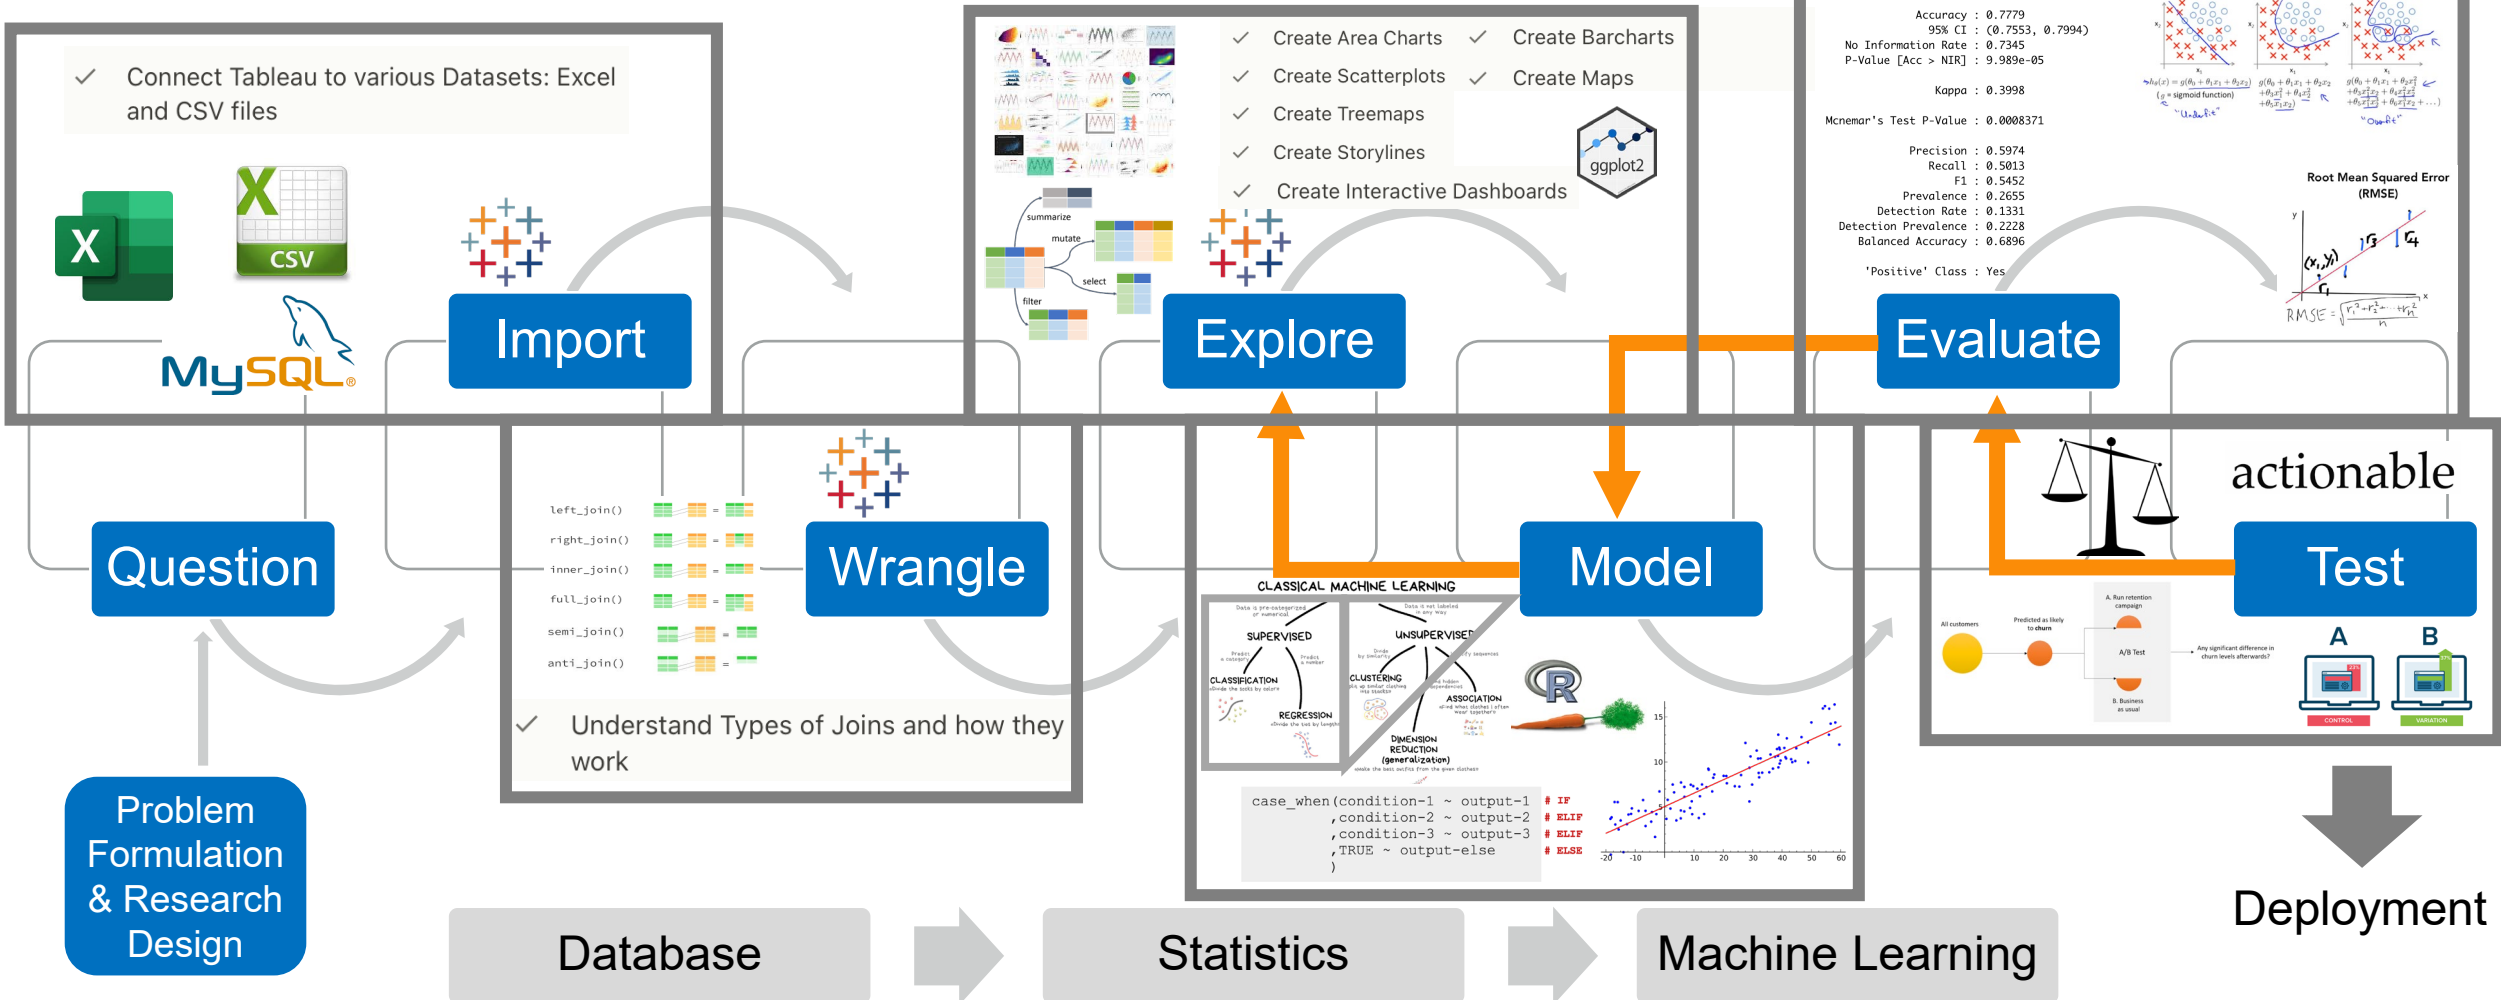
Business Intelligence and Analytics



Terry College of Business
UNIVERSITY OF GEORGIA



Where We Are

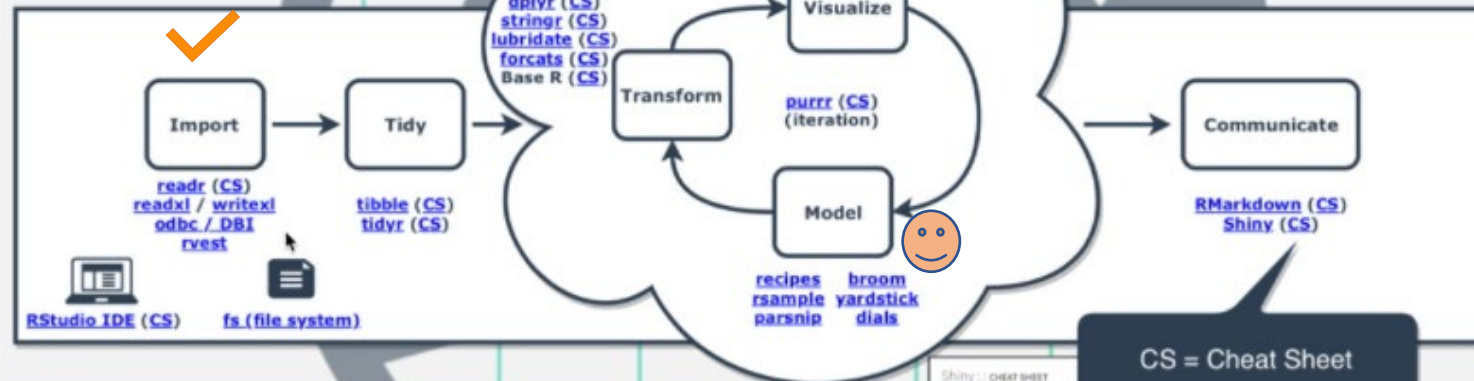


Data Science with R Workflow

The Data Science With R Workflow is available in the book: [R For Data Science](#). If you want to learn R and this workflow for business analysis, take the [R For Business Analysis \(DS4B 101-R\) course](#) through Business Science University.



Click the links for Documentation



Important Resources

- R For Data Science Book: <http://r4ds.had.co.nz/>
- Rmarkdown Book: <https://bookdown.org/yihui/rmarkdown/>
- Data Visualization Book: <https://rkabacoff.github.io/datavis/>
- More Cheatsheets: <https://www.rstudio.com/resources/cheatsheets/>
- tidyverse packages: <https://www.tidyverse.org/>
- Connecting to databases: <https://db.rstudio.com/>
- RMarkdown website: <https://rmarkdown.rstudio.com/>
- Shiny web applications website: <http://shiny.rstudio.com/>
- Jenny Bryan's purrr tutorial: <https://jennybryan.org/>

"Data Science Courses for Business"



Business Science University
university.business-science.io



Import Reminder

Used the **readr** package to import CSV files & the **readxl** package to import Excel files.

```
# Import data from CSV
library(readr)
CoffeeChain <- read_csv("CoffeeChain.csv")

# Import data from Excel
library(readxl)
CoffeeChain <- read_excel("CoffeeChain.xlsx")

# Save data as CSV
write_csv(CoffeeChain, "CoffeeChain.csv")
```



Variable Types

- Numeric
- String
- Date-times
- Binary
- Factors

.....



Terry College of Business
UNIVERSITY OF GEORGIA

Data
Manipulation I

Business Intelligence
Carolina A. de Lima Salge



Tibbles

- A rethinking of data frames, which are essentially data tables.
- One of the unifying features of Tidyverse.
- Tibbles are data frames, but they tweak some older behaviors to make life a little easier.



Tibbles

```
# Load tidyverse package
library(tidyverse)

# Import data as a tibble - using read_csv() function
orderDetails <- read_csv("Order Details.csv")

# Print the tibble
orderDetails

# Import the same data as a data frame - using read.csv() function
orderDetails2 <- read.csv("Order Details.csv")

# Print the data frame
orderDetails2
```



Quick Exercise

Open R Studio, create an R Markdown file, then copy and paste the code from the previous slide into a chunk of R code and execute line by line (Ctrl + Enter / Cmd + Enter)

Do you see something like the figure below? If yes, is it a tibble or a data frame?

```
  `Order ID` Amount Profit Quantity Category  `Sub-Category`  
  <chr>      <dbl> <dbl>    <dbl> <chr>      <chr>  
1 B-25601    1275  -1148      7 Furniture Bookcases  
2 B-25601     66   -12      5 Clothing  Stole  
3 B-25601      8    -2      3 Clothing  Hankerchief  
4 B-25601     80   -56      4 Electronics Electronic Games  
5 B-25602    168  -111      2 Electronics Phones  
6 B-25602    424  -272      5 Electronics Phones  
7 B-25602   2617  1151      4 Electronics Phones  
8 B-25602    561   212      3 Clothing  Saree  
9 B-25602    119    -5      8 Clothing  Saree  
10 B-25603   1355  -60      5 Clothing  Trousers  
# i 1,490 more rows  
# i Use `print(n = ...)` to see more rows
```


Tibbles

```
# To create a tibble from individual
vectors, use tibble()
library(tibble)
gender <- c("m","f","f")
age <- c(5,8,3)
# Create tibble
tb <- tibble(`1 gender`=gender, `2 age`=age)
tb

# From dataframe to tibble
iris_tibble <- as_tibble(iris)
# From tibble to dataframe
iris_df <- as.data.frame(iris_tibble)
```

The parenthesis around the code will also print the tibble

```
(tb <- tibble(x = 1:5,
              y = 1,
              z = x ^ 2 + y))
```

A tibble: 5 x 3

	x	y	z
	<int>	<dbl>	<dbl>
1	1	1	2
2	2	1	5
3	3	1	10
4	4	1	17
5	5	1	26

Tibbles

```
# To subset a tibble use $ or []  
tb$x
```

```
[1] 1 2 3 4 5
```

```
tb[["x"]]
```

```
[1] 1 2 3 4 5
```

```
tb[[1]]
```

```
[1] 1 2 3 4 5
```

```
# tb$x == tb[["x"]] == tb[[1]]
```

```
# $ and [] reference a column in a table (table$column; table[[column]])
```

```
# [] can also be used to reference a specific cell
```

```
# Subset the table to show the second element of the first column (table[[cell2, column1]])
```

```
tb[[2,"x"]]
```

```
[1] 2
```

```
tb[[2,1]]
```

```
[1] 2
```



Dplyr

One of the packages in the Tidyverse that enables the transformation of data

- Look at a subset of the rows—**filter()**
- Reorder rows—**arrange()**
- Rename variables—**rename()**
- Create new variables—**mutate()**
- Collapse values down to a summary—**summarize()**



Data Import

```
library(tidyverse)

orderList <- read_csv("List of Orders.csv")
orderDetails <- read_csv("Order Details.csv")
salesTarget <- read_csv("Sales target.csv")
```



Logical operations

Logical operator	Symbol
EQUAL	==
AND	&
OR	
NOT	!



Missing values

- Missing values are indicated by NA (not available)
- Arithmetic expressions and functions containing missing values generate missing values

```
sum(c(1, NA, 2))
```

- Use the na.rm=T option to exclude missing values for calculations

```
sum(c(1, NA, 2), na.rm=TRUE)
```



The Pipe

The `%>%` focuses on the transformations, not what is being transformed, which makes the code easier to read.

- Only order details with products in Category == “Furniture”
- Take the `orderDetails` dataset **then** `filter()`

```
orderDetails %>%  
  filter(., Category == "Furniture")
```

Filter Rows

A reference to
orderDetails

```
orderDetails %>%  
  filter(., Category == "Furniture")  
# A tibble: 243 x 6  
  `Order ID` Amount Profit Quantity Category `Sub-Category`  
  <chr>      <dbl> <dbl>    <dbl> <chr>      <chr>  
1 B-25601    1275  -1148     7 Furniture Bookcases  
2 B-25603     24   -30     1 Furniture Chairs  
3 B-25608   1364 -1864     5 Furniture Tables  
4 B-25608    476    0     3 Furniture Chairs  
5 B-25610     30   -5     2 Furniture Furnishings  
6 B-25612    259  -55     2 Furniture Chairs  
7 B-25614    494   54     4 Furniture Bookcases  
8 B-25618    362  127     1 Furniture Bookcases  
9 B-25626   1103  -276     3 Furniture Chairs  
10 B-25628     35   -8     2 Furniture Furnishings  
# ... with 233 more rows
```

Only order details with
products in
Category == "Furniture"



Filter Rows

and



```
orderDetails %>%  
  filter(., Category == "Furniture", Quantity > 1)  
# A tibble: 223 x 6  
  `Order ID` Amount Profit Quantity Category `Sub-Category`  
  <chr>      <dbl> <dbl>    <dbl> <chr>      <chr>  
1 B-25601    1275  -1148     7 Furniture Bookcases  
2 B-25608    1364  -1864     5 Furniture Tables  
3 B-25608     476     0     3 Furniture Chairs  
4 B-25610      30     -5     2 Furniture Furnishings  
5 B-25612     259    -55     2 Furniture Chairs  
6 B-25614     494     54     4 Furniture Bookcases  
7 B-25626    1103   -276     3 Furniture Chairs  
8 B-25628      35     -8     2 Furniture Furnishings  
9 B-25631      89    -89     2 Furniture Furnishings  
10 B-25634     389    -83     3 Furniture Chairs  
# ... with 213 more rows
```

Only order details with
products in
Category == "Furniture"
& Quantity > 1



Filter Rows

```
orderDetails %>%  
  filter(., Category == "Furniture" & Quantity > 1)  
# A tibble: 223 x 6  
  `Order ID` Amount Profit Quantity Category `Sub-Category`  
  <chr>      <dbl> <dbl>    <dbl> <chr>      <chr>  
1 B-25601    1275 -1148      7 Furniture Bookcases  
2 B-25608    1364 -1864      5 Furniture Tables  
3 B-25608     476    0      3 Furniture Chairs  
4 B-25610      30   -5      2 Furniture Furnishings  
5 B-25612     259  -55      2 Furniture Chairs  
6 B-25614     494   54      4 Furniture Bookcases  
7 B-25626    1103 -276      3 Furniture Chairs  
8 B-25628      35   -8      2 Furniture Furnishings  
9 B-25631      89  -89      2 Furniture Furnishings  
10 B-25634     389  -83      3 Furniture Chairs  
# ... with 213 more rows
```

and



Only order details with
products in
Category == "Furniture"
& Quantity > 1



Filter Rows

```
orderDetails %>%  
  filter(., Category == "Furniture" | Quantity > 1)
```

```
# A tibble: 1,388 x 6
```

	`Order ID`	Amount	Profit	Quantity	Category	`Sub-Category`
	<chr>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
1	B-25601	1275	-1148	7	Furniture	Bookcases
2	B-25601	66	-12	5	Clothing	Stole
3	B-25601	8	-2	3	Clothing	Hankerchief
4	B-25601	80	-56	4	Electronics	Electronic Games
5	B-25602	168	-111	2	Electronics	Phones
6	B-25602	424	-272	5	Electronics	Phones
7	B-25602	2617	1151	4	Electronics	Phones
8	B-25602	561	212	3	Clothing	Saree
9	B-25602	119	-5	8	Clothing	Saree
10	B-25603	1355	-60	5	Clothing	Trousers

```
# ... with 1,378 more rows
```

or



Only order details with
products in
Category == "Furniture"
or Quantity > 1



Filter Rows

```
orderDetails %>%  
  filter(., Category %in% c("Furniture", "Clothing"))
```

```
# A tibble: 1,192 x 6
```

	`Order ID` <chr>	Amount <dbl>	Profit <dbl>	Quantity <dbl>	Category <chr>	`Sub-Category` <chr>
1	B-25601	1275	-1148	7	Furniture	Bookcases
2	B-25601	66	-12	5	Clothing	Stole
3	B-25601	8	-2	3	Clothing	Hankerchief
4	B-25602	561	212	3	Clothing	Saree
5	B-25602	119	-5	8	Clothing	Saree
6	B-25603	1355	-60	5	Clothing	Trousers
7	B-25603	24	-30	1	Furniture	Chairs
8	B-25603	193	-166	3	Clothing	Saree
9	B-25603	180	5	3	Clothing	Trousers
10	B-25603	116	16	4	Clothing	Stole

```
# ... with 1,182 more rows
```

Only order details with
products in
Category in “Furniture”
or “Clothing”



Practice Question

Mike is trying to create a new table called **posBAprofit** that filters the orderDetails table to show observations where profit is above zero but below the average. How could you create this table in R?



Reorder (Arrange) Rows

```
orderDetails %>%  
  arrange(., desc(Profit))  
# A tibble: 1,500 x 6  
  `Order ID` Amount Profit Quantity Category `Sub-Category`  
  <chr>      <dbl>  <dbl>    <dbl> <chr>      <chr>  
1 B-25973    4141   1698     13 Electronics Printers  
2 B-25602    2617   1151      4 Electronics Phones  
3 B-25761    2188   1050      5 Furniture Bookcases  
4 B-25923    3873    891      6 Electronics Phones  
5 B-25830    1954    782      3 Electronics Phones  
6 B-26073    1514    742      4 Electronics Printers  
7 B-25853    2093    721      5 Furniture Chairs  
8 B-26093    2847    712      8 Electronics Printers  
9 B-25862    2061    701      5 Furniture Bookcases  
10 B-25656   1389    680      7 Clothing Saree  
# ... with 1,490 more rows
```



Selecting Variables

```
orderDetails %>%  
  select(., c(Amount:Category))  
# A tibble: 1,500 × 4  
  Amount Profit Quantity Category  
  <dbl>   <dbl>   <dbl> <chr>  
1    1275   -1148       7 Furniture  
2      66    -12       5 Clothing  
3       8     -2       3 Clothing  
4      80   -56       4 Electronics  
5     168  -111       2 Electronics  
6     424  -272       5 Electronics  
7    2617   1151       4 Electronics  
8     561    212       3 Clothing  
9     119     -5       8 Clothing  
10    1355   -60       5 Clothing  
# ... with 1,490 more rows
```



Rename Variables

orderDetails %>%
 rename(., profit = Profit)

New name Old name

A tibble: 1,500 x 6

	`Order ID`	Amount	profit	Quantity	Category	`Sub-Category`
	<chr>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
1	B-25601	1275	-1148	7	Furniture	Bookcases
2	B-25601	66	-12	5	Clothing	Stole
3	B-25601	8	-2	3	Clothing	Hankerchief
4	B-25601	80	-56	4	Electronics	Electronic Games
5	B-25602	168	-111	2	Electronics	Phones
6	B-25602	424	-272	5	Electronics	Phones
7	B-25602	2617	1151	4	Electronics	Phones
8	B-25602	561	212	3	Clothing	Saree
9	B-25602	119	-5	8	Clothing	Saree
10	B-25603	1355	-60	5	Clothing	Trousers

... with 1,490 more rows

Add New Variables

Name of the new variable

Value of the new variable

```
orderDetails %>%  
  mutate(ProfitN = (Profit - min(Profit)) / (max(Profit) - min(Profit)))
```

A tibble: 1,500 x 7

	`Order ID` <chr>	Amount <dbl>	Profit <dbl>	Quantity <dbl>	Category <chr>	`Sub-Category` <chr>	ProfitN <dbl>
1	B-25601	1275	-1148	7	Furniture	Bookcases	0.226
2	B-25601	66	-12	5	Clothing	Stole	0.535
3	B-25601	8	-2	3	Clothing	Hankerchief	0.538
4	B-25601	80	-56	4	Electronics	Electronic Games	0.523
5	B-25602	168	-111	2	Electronics	Phones	0.508
6	B-25602	424	-272	5	Electronics	Phones	0.465
7	B-25602	2617	1151	4	Electronics	Phones	0.851
8	B-25602	561	212	3	Clothing	Saree	0.596
9	B-25602	119	-5	8	Clothing	Saree	0.537
10	B-25603	1355	-60	5	Clothing	Trousers	0.522

... with 1,490 more rows



Add New Variables

Name of the new variable

```
orderDetails %>%  
  mutate(., ProfitClass = case_when(Profit > 0 ~ "Positive",  
    Profit < 0 ~ "Negative", Profit == 0 ~ "Zero"))
```

A tibble: 1,500 x 7

	`Order ID`	Amount	Profit	Quantity	Category	`Sub-Category`	ProfitClass
	<chr>	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<chr>
1	B-25601	1275	-1148	7	Furniture	Bookcases	Negative
2	B-25601	66	-12	5	Clothing	Stole	Negative
3	B-25601	8	-2	3	Clothing	Hankerchief	Negative
4	B-25601	80	-56	4	Electronics	Electronic Games	Negative
5	B-25602	168	-111	2	Electronics	Phones	Negative
6	B-25602	424	-272	5	Electronics	Phones	Negative
7	B-25602	2617	1151	4	Electronics	Phones	Positive
8	B-25602	561	212	3	Clothing	Saree	Positive
9	B-25602	119	-5	8	Clothing	Saree	Negative
10	B-25603	1355	-60	5	Clothing	Trousers	Negative

... with 1,490 more rows



Grouped Summaries

Calculate the average profit for each “Category” and “Sub-Category” combination. Then, sort the results descendingly to identify the most profitable groups.

Take the `orderDetails` dataset **then** `group_by()` **then** `summarise()` **then** `arrange()`

```
orderDetails %>%  
  group_by(Category, `Sub-Category`) %>%  
  summarize(`Average Profit` = mean(Profit, na.rm = TRUE)) %>%  
  arrange(desc(`Average Profit`))
```

Grouped Summaries

Grouped
variables

Summarized function
and variable

Name of the new
summarized
variable

```
orderDetails %>%
  group_by(Category, `Sub-Category`) %>%
  summarize(`Average Profit` = mean(Profit, na.rm = TRUE)) %>%
  arrange(desc(`Average Profit`))
```

A tibble: 17 x 3

Groups: Category [3]

	Category	`Sub-Category`	`Average Profit`
	<chr>	<chr>	<dbl>
1	Electronics	Printers	80.6
2	Clothing	Trousers	73
3	Furniture	Bookcases	61.9
4	Electronics	Accessories	49.4
5	Electronics	Phones	26.6
6	Clothing	T-shirt	19.5
7	Clothing	Shirt	16.4
8	Clothing	Stole	13.3
9	Furniture	Furnishings	11.6
10	Clothing	Hankerchief	10.6
11	Furniture	Chairs	7.80
12	Clothing	Leggings	4.91
13	Clothing	Kurti	3.85
14	Clothing	Skirt	3.67
15	Clothing	Saree	1.68
16	Electronics	Electronic Games	-15.6
17	Furniture	Tables	-236.



Useful Summary Functions

- Location
 - **mean(x)** and **median(x)**
- Spread
 - **sd(x)** and **IQR(x)**
- Rank
 - **min(x)**, **quantile(x, 0.25)**, and **max(x)**
- Count
 - **n(x)**, **sum(!is.na(x))**, and **n_distinct(x)**



dplyr helper functions

- Use with select()
 - starts_with()
 - matches column names that begin with certain characters
 - ends_with()
 - matches column names that end with specific characters
 - contains()
 - matches column names that contain a sequence of characters



Try Yourself!

Using the *orderList* dataset, create a table with the number of orders per state and sort it descendingly.

